

Vishal Kumar Gowda

Bengaluru, Karnataka, India | +91 8088833875 | vishalgowdask65@gmail.com | [LinkedIn](#) | [GitHub](#) | [LeetCode](#) | [Portfolio](#)

EDUCATION

Ajeenkya D Y Patil University (Newton School Of Technology)

2025-2029

B.Tech in Computer Science and Engineering, Specialisation in AI/ML — CGPA: 7.32/10

EXPERIENCE

AI/ML Intern, Smarta Tech

June 2026 – Present

- Designed and implemented a concurrent, asynchronous request-handling architecture for calls to **LLM APIs**, cutting pipeline latency and improving throughput for production inference workloads.
- Independently built and shipped a custom production-grade **RAG application** with a voice assistant feature as an internship deliverable, engineering it entirely on free/open-source tools with zero paid subscriptions as per condition.
- Architected a retrieval system integrating **vector databases** with LLM APIs to improve context relevance and accuracy of generated responses for downstream applications.
- Owned a production ML pipeline **end-to-end** — data collection, data mining, preprocessing, model development, and evaluation — for a real-world AI application, taking it from prototype to deployment.
- Collaborated directly with stakeholders to translate business requirements into data-driven **web solutions**, using systematic feature engineering.

Founder & Lead, MediConnect (AIStartup)

2025 –

2026

- Led** a 5-person founding team building a healthcare connectivity platform, owning workflow, planning, milestone management, and risk preparedness.
- I was **responsible** for workflow planning, milestone management, **team management**, and analysing potential problems in the feature and preparing the team well in advance.
- Competed in the **Shark Pool Hackathon**, an investor-pitch competition; **cleared 3 rounds** to reach the final investors' round.

PROJECTS

Custom RAG Application

June 2026 – Present

[GitHub](#) | [Demo Video](#)

- Built a custom RAG application from **scratch** for a real-world use case, featuring PDF upload with an **end-to-end** document Q&A pipeline, maintained with **production-level code architecture**.
- Shipped iterative versioned releases (v1, v2, ...), adding a new feature with every update.
- Engineered a **voice assistant feature** and ran the entire stack on free/open-source tools with zero paid subscriptions.

House Price Prediction

June 2026

[GitHub](#) | [Live Demo](#)

- Built an end-to-end ML system in Python covering data preprocessing, feature engineering, model training, and deployment.
- Benchmarked **Linear Regression**, Random Forest, and XGBoost; XGBoost performed best ($R^2 \approx 0.41$, RMSE $\approx ₹1.71L$), outperforming the linear baseline by capturing non-linear feature interactions.
- Deployed a Streamlit application enabling real-time property price predictions for end users.

Smart Store

April 2026

[GitHub](#) | [Live Demo](#)

- Built a **responsive** e-commerce frontend in React with authentication, protected routes, and form validation.
- Designed global state management with **Context API** and **reusable components** to optimise UI performance and maintainability; full-stack and SQL database in progress.

TECHNICAL SKILLS

Core CS: Data Structures & Algorithms, Object-Oriented Design, Concurrency, Clean Architecture, SOLID Principles, Production Codebase Architecture & Asynchronous Programming, Software Design, problem-solving, active learning

Languages: Python, C++, JavaScript, HTML, CSS

Frameworks & Libraries: React, Node.js, Express.js, FastAPI, LangChain, Pytest, Pandas, NumPy, OpenCV, Matplotlib, Seaborn, TensorFlow

Cloud & DevOps: Azure, Docker, Git/GitHub, GitHub Actions, Virtualisation

AI / ML: Machine Learning, LLMs, Retrieval-Augmented Generation (RAG) Pipelines, Hugging Face, OpenAI API, MCP, Prompt Engineering, Vector Databases, ChromaDB, Embeddings, Fine-tuning, Ollama

APIs & Tools: REST APIs, OpenAPI (Swagger)

Testing: Unit Testing, Integration Testing

CERTIFICATIONS

Claude Code With The Anthropic API — Anthropic

March 2026

- Implemented RAG pipelines, MCP integrations, and advanced prompting techniques using the Anthropic API to build context-aware AI applications.**